

Computing Summary Statistics for Daily Data

Dave Lorenz

April 18, 2014

These examples demonstrate how to compute selected summary statistics for daily stream-flow data. The examples can easily be extended to other statistics or data types.

```
> # Load the USGSwsGraphs package
> library(USGSwsBase)
> # Retrieve streamflow data for the Raccoon River at Van Meter, IA
> # for the 30-year period beginning 1980-10-01.
> # Use the renCol function to rename the strflow column to Flow
> RRVM <- renCol(readNWIS("05484500", begin.date="1980-10-01",
+                          end.date="2010-09-30"))
> # Print the first and last few rows of the data
> head(RRVM)
```

	agency_cd	site_no	datetime	Flow	Flow_cd
1	USGS	05484500	1980-10-01	278	A
2	USGS	05484500	1980-10-02	258	A
3	USGS	05484500	1980-10-03	235	A
4	USGS	05484500	1980-10-04	225	A
5	USGS	05484500	1980-10-05	221	A
6	USGS	05484500	1980-10-06	216	A

```
> tail(RRVM)
```

	agency_cd	site_no	datetime	Flow	Flow_cd
10952	USGS	05484500	2010-09-25	4180	A
10953	USGS	05484500	2010-09-26	6150	A
10954	USGS	05484500	2010-09-27	7060	A
10955	USGS	05484500	2010-09-28	6960	A
10956	USGS	05484500	2010-09-29	6240	A
10957	USGS	05484500	2010-09-30	5380	A

```
> # Check for missing values
> with(RRVM, screenData(datetime, Flow, year = "water"))
```

No missing data between 1980-10-01 and 2010-09-30

1 Computing Daily Mean Values

The simplest and most straightforward way to compute summary statistics from arbitrarily grouped data is to use the `tapply` function. At its simplest, it requires only three arguments—`X`, the data to summarize; `INDEX`, the grouping data; and `FUN`, the summary statistic function.

The `USGSwsBase` package contains the `baseDay` function that can be used to group data by day, so that all data for each day, including February 29, can be summarized. The output can be arranged so that the sequence represents the calendar-, water- or climate-year; beginning January 1, October 1, or April 1.

The following script demonstrates how to use the `tapply` and `baseDay` functions to compute the daily mean streamflow for the previously retrieved data. It uses the `with` function to facilitate referring to columns in the dataset.

```
> # There are no missing values, so only need the basic
> # 3 arguments for tapply
> RRVM.daily <- with(RRVM, tapply(Flow,
+   baseDay(datetime, numeric=FALSE, year="water"), mean))
> # Print the first and last few values of the output
> head(RRVM.daily)

    Oct 01    Oct 02    Oct 03    Oct 04    Oct 05    Oct 06
804.3000 777.7000 792.7667 884.4667 976.8333 984.5000

> tail(RRVM.daily)

    Sep 25    Sep 26    Sep 27    Sep 28    Sep 29    Sep 30
1059.5667 1093.9333 1097.7333 1067.0000 1060.5667  988.2333
```

The output from `tapply` is an array. Because the output from this example is an array of one dimension, it is printed in the form of a named vector. Had the summary statistic function been `quantile`, for example, the output would have been a list.

The `tapply` function is very powerful and easy to use. But there are times when we want the output in the form of a dataset rather than a vector or array. In those cases, the `aggregate` function is a better alternative. The `aggregate` function has several usage options. The script below demonstrates how to build a formula to compute the same statistics that we computed in the previous script. Early versions of `aggregate` required the output of the summary statistic function to be a scalar, but that is no longer a limitation.

```
> # There are no missing values
> RRVM.dailyDF <- aggregate(Flow ~
+   baseDay(datetime, numeric=FALSE, year="water"),
+   data=RRVM, FUN=mean)
> # Print the first and last few values of the output
> head(RRVM.dailyDF)
```

```

      baseDay(datetime, numeric = FALSE, year = "water")      Flow
1          Oct 01 804.3000
2          Oct 02 777.7000
3          Oct 03 792.7667
4          Oct 04 884.4667
5          Oct 05 976.8333
6          Oct 06 984.5000

> tail(RRVM.dailyDF)

      baseDay(datetime, numeric = FALSE, year = "water")      Flow
361          Sep 25 1059.5667
362          Sep 26 1093.9333
363          Sep 27 1097.7333
364          Sep 28 1067.0000
365          Sep 29 1060.5667
366          Sep 30  988.2333

> # Change the name of the grouping column
> names(RRVM.dailyDF)[1] <- "Day"

```

Note that the grouping column, now called Day, is a factor. There are times when it would be better to be a simple character. It can easily be converted by the executing the expression: `RRVM.dailyDF$Day <- as.character(RRVM.dailyDF$Day)`.

2 Computing Annual Mean Values

The example above can easily be expanded to any grouping that the user desires. This example computes annual means by water year. The `waterYear` function in `USGSwsBase` is used to group the data by water year.

```
> # There are no missing values
> RRVM.wyDF <- aggregate(Flow ~
+   waterYear(datetime),
+   data=RRVM, FUN=mean)
> # Change the name of the grouping column
> names(RRVM.wyDF)[1] <- "WaterYear"
> # Print the first few values of the output
> head(RRVM.wyDF)
```

	WaterYear	Flow
1	1981	373.6986
2	1982	2198.0192
3	1983	4572.0767
4	1984	4005.3962
5	1985	1009.1425
6	1986	2659.5178

Other grouping functions include `year` (calendar year) in `lubridate`, `month` (month) in `lubridate`, `seasons` (user-defined seasons) in `USGSwsBase`. Refer to the documentation for each of these function for a description of the arguments.

3 Computing Year and Month Mean Values

Aggregation can also be done by multiple grouping variables. This example computes the mean streamflow for each month by year. This example uses the `year` and the `month` functions because the output is sorted by groups. The sequence of the groups in the call is important—the sorting is done in the order specified in the formula. For this example, the data are sorted by month and then by year, which in this case, keeps the order correct; grouping by water year would misplace October, November and December. For a calendar year table, the months are in the correct order.

```
> # There are no missing values
> RRVM.my <- aggregate(Flow ~ month(datetime, label=TRUE) + year(datetime),
+   data=RRVM, FUN=mean)
> # Rename columns 1 and 2
> names(RRVM.my)[1:2] <- c("Month", "Year")
> # Print the first few values of the output
> head(RRVM.my)
```

	Month	Year	Flow
1	Oct	1980	226.0323
2	Nov	1980	201.0333
3	Dec	1980	137.0000
4	Jan	1981	102.0000
5	Feb	1981	256.9643
6	Mar	1981	166.1613

The output dataset may be used as is, or it could be restructured to a table of monthly values for each water year. To do that, a column of water year must be added and the order months must be set to correspond to those in the water year. The following script uses the `%in%` function to recode the water year from the year.

```
> # Create new object, compute the water year and round Flow
> RRVM.myTbl <- transform(RRVM.my,
+   WY=ifelse(Month %in% c("Oct", "Nov", "Dec"), Year + 1, Year),
+   Flow=signif(Flow, 3))
> # Reorder months
> RRVM.myTbl$Month <- factor(RRVM.myTbl$Month,
+   levels=c("Oct", "Nov", "Dec", "Jan", "Feb", "Mar", "Apr", "May",
+   "Jun", "Jul", "Aug", "Sep"))
> # Restructure the dataset, overwrite the new object
> RRVM.myTbl <- group2row(RRVM.myTbl, "WY", "Month", "Flow")
> # Print the first few values of the output, set width for Vignette
> options(width=70)
> head(RRVM.myTbl)
```

	WY	Oct.Flow	Nov.Flow	Dec.Flow	Jan.Flow	Feb.Flow	Mar.Flow	Apr.Flow
1	1981	226	201	137	102	257	166	207

2	1982	332	409	509	221	2180	4010	2380
3	1983	2450	2250	3090	2520	4970	8320	10600
4	1984	1230	3060	1760	865	5440	3260	9100
5	1985	558	776	1130	1390	1310	1280	1820
6	1986	436	268	247	278	474	3530	3230
	May.Flow	Jun.Flow	Jul.Flow	Aug.Flow	Sep.Flow			
1	360	550	1170	727	364			
2	6650	3960	3860	913	913			
3	6300	6090	6920	853	551			
4	9260	11000	2530	663	274			
5	1900	1070	483	192	223			
6	7240	4470	6730	2590	2200			

Note that this example used the `group2row` function in `USGSwsBase`. The `reshape` function in `stats` and `stack` and `unstack` functions in `utils` are other functions that will restructure data.